

REPORT PROGRESS 2

Main task: Set up real-time data streaming from database

Tool: Redis, Spark

Task description:

- Upload data from Redis to Spark
- Set up real-time data streaming

Task execution process

Step 1: Import library

Import the libraries needed for uploading data

```
import time
from pyspark.sql import SparkSession
import redis
from pyspark import SparkContext
```

Step 2: Create Spark Session and Connect Redis

```
# Create Spark Session
spark = SparkSession.builder.appName("RedisToSparkStreaming").getOrCreate()
sc = spark.sparkContext

# Connect Redis
redis_host = "192.168.126.131"
redis_port = 6379
batch_size = 10000 # Number of keys per batch
```

Step 3: Get Redis keys

```
def get_redis_keys():
    "Use SCAN to get a list of keys from Redis"
    r = redis.StrictRedis(host=redis_host, port=redis_port, decode_responses=True)
    cursor = 0
    keys = []
    while True:
        cursor, batch_keys = r.scan(cursor, count=batch_size)
        keys.extend(batch_keys)
        if cursor == 0: # Out of keys in Redis
            break
    print(f"Number of keys currently in Redis: {len(keys)}")
    return keys
```

- Create a Redis connection using StrictRedis with information from redis_host and redis_port.
- decode_responses=True: Makes the data returned from Redis be strings instead of bytes.
- Initialize variables.
- Use scan() to get a batch of keys from Redis.
- cursor: Redis uses cursor to iterate through the data piece by piece.
- If cursor=0, it means start from the beginning.
- If scan() returns cursor=0 after one iteration, it means all keys have been scanned.
- Save the retrieved keys to the keys list.
- When cursor == 0, it means there is no more data to scan, exit the loop.
- Return the keys list.

Step 4: Fetches data from Redis

```
def fetch_redis_data(keys_batch):
    "Fetch data from Redis with smaller batch and check data type before GET"
    r = redis.StrictRedis(host=redis_host, port=redis_port, decode_responses=True)

    batch_size = 1000 # Split batches to avoid overloading
    data = []
    keys_batch = list(keys_batch)
    for i in range(0, len(keys_batch), batch_size):
        pipe = r.pipeline()
        sub_batch = keys_batch[i : i + batch_size] # Split batch

        # Check data type before GET
        valid_keys = [key for key in sub_batch if r.type(key) == "string"]
        for key in valid_keys:
            pipe.get(key) # GET only keys of string type

        values = pipe.execute()

        for key, value in zip(valid_keys, values):
            if value is not None:
                try:
                    user_id, anime_id = key.split("_")
                    rating = float(value)
                    data.append((int(user_id), int(anime_id), rating))
                except ValueError:
                    continue # Bỏ qua key lỗi
    return data
```

- This function fetches data from Redis while optimizing performance by using batch processing.

- A new Redis connection `r` is established using `StrictRedis`, with `decode_responses=True` to ensure values are returned as strings instead of bytes.
- `batch_size = 1000`: The function processes keys in chunks of 1000 to prevent overwhelming Redis.
- `data = []`: A list to store the final processed results.
- `keys_batch = list(keys_batch)`: Ensures that `keys_batch` is a list, in case it was provided as another iterable (e.g., a generator).
- Looping through the keys in chunks of `batch_size` to fetch data in smaller portions.
- `r.pipeline()`: A Redis pipeline is used to reduce the number of network requests, improving efficiency.
- `sub_batch = keys_batch[i : i + batch_size]`: Extracts a smaller subset of keys for processing.
- Filters out keys that are not of type string before calling `GET`.
- This prevents unnecessary operations on incompatible key types (e.g., lists, sets, hashes).
- Adds `GET` commands for all valid keys to the Redis pipeline.
- Executes all `GET` commands in one batch to minimize network overhead.
- Iterates through the fetched values and their corresponding keys.
- Parses the key into `user_id` and `anime_id` (assuming they are separated by `_`).
- Converts the value to float since it represents a rating.
- Appends the parsed data as a tuple (`user_id`, `anime_id`, `rating`) to the `data` list.
- If an error occurs during conversion (e.g., improperly formatted keys or values), the function skips the key instead of crashing.
- Returns the final list of parsed (`user_id`, `anime_id`, `rating`) tuples.

Step 5: Streaming data

```

def streaming():
    old_redis_keys = 0
    # Streaming simulation loop (updates every 10 seconds)
    while True:
        print(" Fetching data from Redis")
        # Get list of keys from Redis
        redis_keys = get_redis_keys()
        if len(redis_keys) != old_redis_keys:
            # Divide the key into multiple partitions for Spark to process in parallel
            num_partitions = 500 # Split into 500 partitions
            rdd_keys = sc.parallelize(redis_keys, numSlices=num_partitions)

            # Use mapPartitions to reduce Redis connection times
            rdd_data = rdd_keys.mapPartitions(fetch_redis_data)

            # Convert to DataFrame
            df = spark.createDataFrame(rdd_data, ["user_id", "anime_id", "rating"])

            # Show DataFrame
            show_df(df)
            old_redis_keys = len(redis_keys)
        else:
            print("The data set has no changes")
            show_df(df)
            old_redis_keys = len(redis_keys)

        print("Wait 10 seconds before retrieving new data...")
        time.sleep(10) # Wait 10 seconds before retrieving next data

```

- The streaming() function is designed to continuously fetch data from Redis, process it using Apache Spark, and display the results every 10 seconds. This simulates a real-time data ingestion pipeline.
- old_redis_keys = 0: This variable keeps track of the number of previously fetched Redis keys. It is used to detect if new data has been added to Redis, avoiding unnecessary processing.
- Runs an infinite loop that fetches data every 10 seconds. Acts as a real-time data ingestion pipeline where Redis is polled for updates.
- Calls get_redis_keys(), which retrieves all keys stored in Redis. The number of keys will determine whether data processing is needed.
- Compares the new key count with old_redis_keys, if the number of keys has changed, new data is available for

processing. This avoids unnecessary computation when the data remains the same.

- Redis keys are divided into 500 partitions for parallel processing using Spark.
- `sc.parallelize(redis_keys, numSlices=500)` creates an RDD (Resilient Distributed Dataset) to distribute data across multiple executors. This helps scale the processing across multiple nodes.
- `mapPartitions(fetch_redis_data)` applies `fetch_redis_data()` to each partition of keys.
- Instead of opening a new Redis connection for each key, it processes keys in batches, reducing the number of Redis calls. Optimized for performance when dealing with large datasets.
- The processed RDD is converted into a Spark DataFrame. The DataFrame has three columns: `user_id`, `anime_id`, and `rating`. This DataFrame is now ready for further analytics or machine learning models.
- Updates `old_redis_keys` to store the new state of Redis keys. This ensures the function only processes newly added data in the next iteration.
- If no new data is found, it skips reprocessing and just displays the existing DataFrame.
- This prevents unnecessary computation, making the function efficient.

Result:

```
df.where(df.user_id == 1).show(20)
```

user_id	anime_id	rating
1	591	6.0
1	400	7.0
1	68	7.0
1	45	8.0
1	167	8.0
1	263	10.0
1	3226	7.0
1	93	8.0
1	849	10.0
1	139	9.0
1	73	10.0
1	44	10.0
1	2001	8.0
1	214	3.0
1	165	8.0
1	267	8.0
1	478	7.0
1	853	9.0
1	7	9.0
1	415	7.0

Show first 20 rows are reviews of user with id 1 for anime in dataframe

```
row_count = df.count()
print(f"Sample row count: {row_count}")
```

```
Sample row count: 24325191
```

Count the number of rows in the current Dataframe